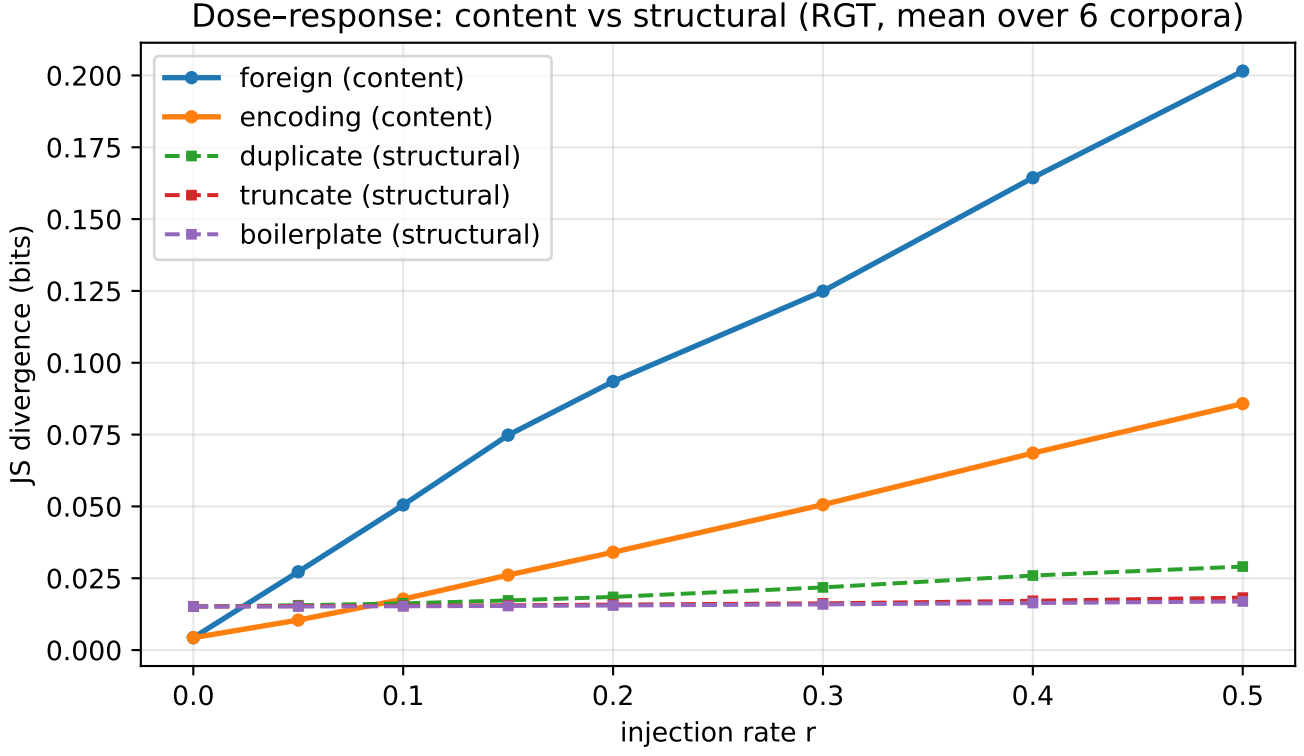
Dose–response: content vs structural (RGT, mean over 6 corpora)
foreign (content)
encoding (content)
duplicate (structural)
truncate (structural)
boilerplate (structural)
JS divergence (bits)
0.200
0.175
0.150
0.125
0.100
0.075
0.050
0.025
0.000
0.0
0.1
0.2
0.3
0.4
0.5
injection rate r

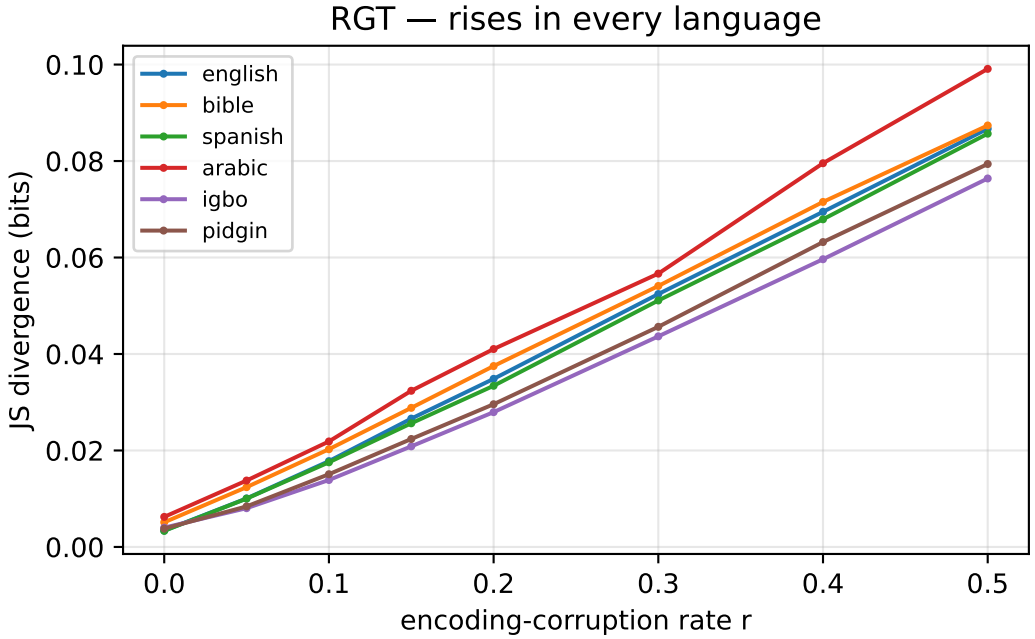
RGT — rises in every language
JS divergence (bits)
encoding-corruption rate r
english
bible
spanish
arabic
igbo
pidgin
0.00
0.02
0.04
0.06
0.08
0.10
0.0
0.1
0.2
0.3
0.4
0.5

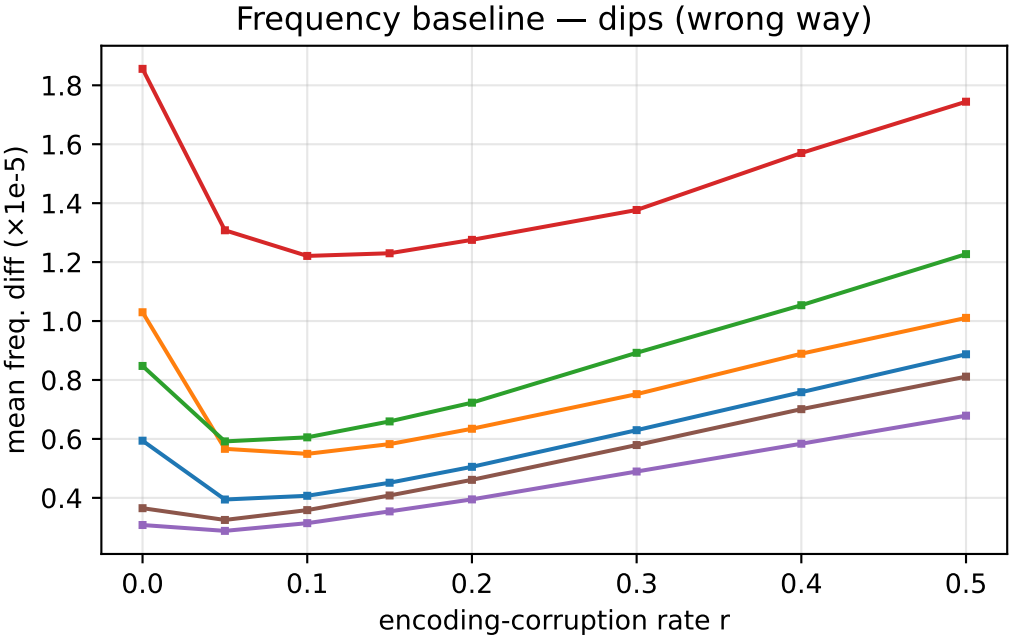
Frequency baseline — dips (wrong way)
mean freq. diff (×1e-5)
encoding-corruption rate r
0.4
0.6
0.8
1.0
1.2
1.4
1.6
1.8
0.0
0.1
0.2
0.3
0.4
0.5

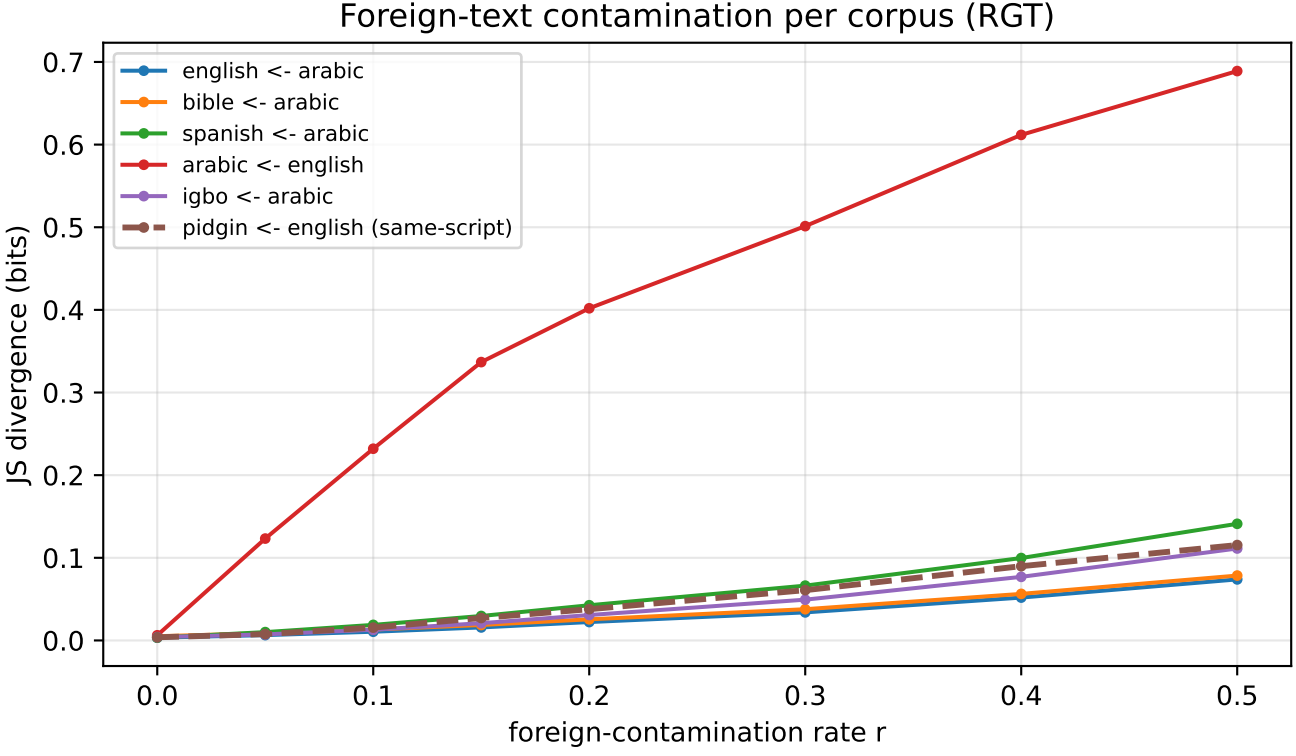
Foreign-text contamination per corpus (RGT)
english <- arabic
bible <- arabic
spanish <- arabic
arabic <- english
igbo <- arabic
pidgin <- english (same-script)
JS divergence (bits)
0.0
0.1
0.2
0.3
0.4
0.5
0.6
0.7
0.0
0.1
0.2
0.3
0.4
0.5
foreign-contamination rate r

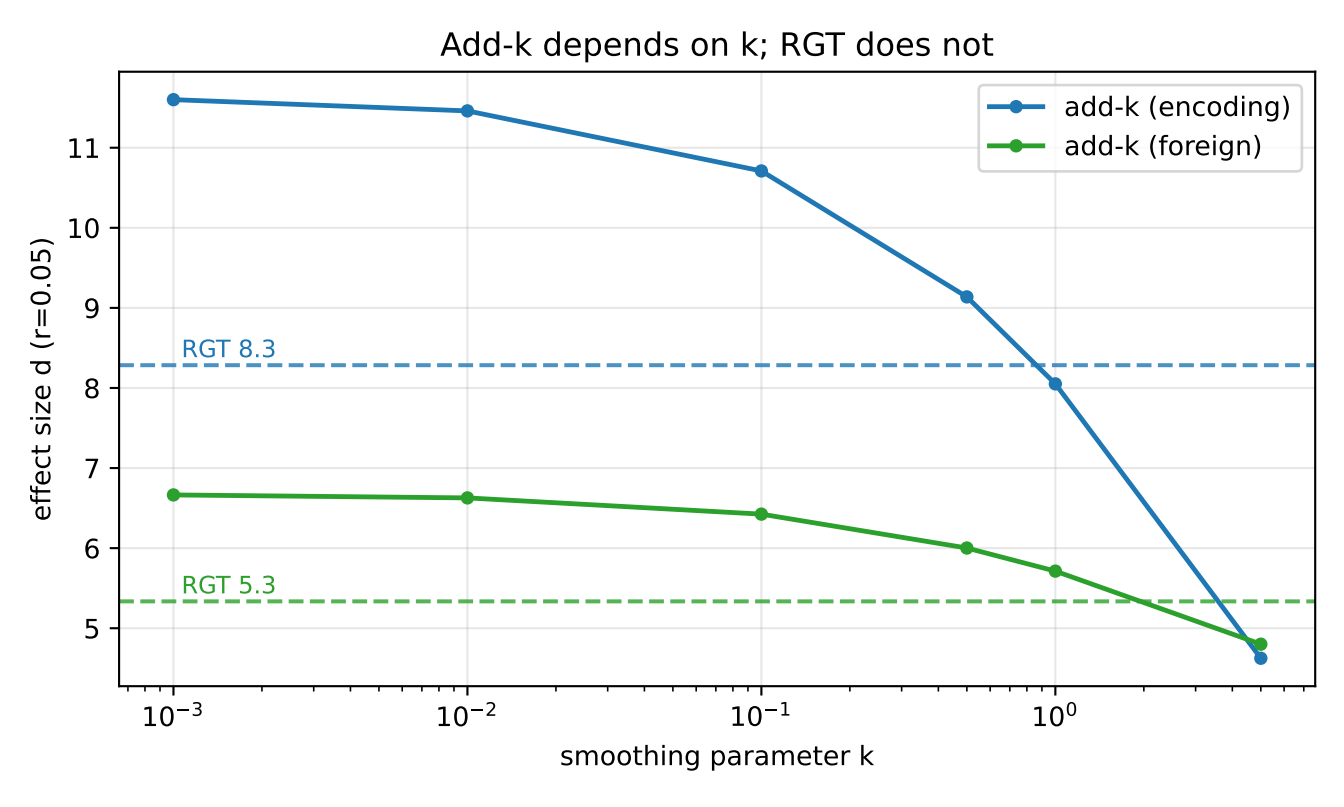
Add-k depends on k; RGT does not
add-k (encoding)
add-k (foreign)
RGT 8.3
RGT 5.3
effect size d (r=0.05)
smoothing parameter k
11
10
9
8
7
6
5
10^{-3}
10^{-2}
10^{-1}
10^{0}